

Hate Speech Detection using Transformers and Deep Learning

By:

Farheen Fatima
Mohamed Derbeli

- Introduction
- Dataset
- Data preprocessing
- Models
- Results

The term hate speech is understood as any type of verbal, written or behavioral communication that attacks or uses derogatory or discriminatory language against a person or group based on what they are. In other words, based on their religion, ethnicity, nationality, race, color, ancestry, sex or other identity factor, In the problem, we will take you through a hate speech detection model with Machine Learning, Deep Learning and Python.

Hate Speech Detection is generally a task of sentiment classification. So, for training a model that can classify hate speech from a certain piece of text can be achieved by training it on data that is generally used to classify sentiments. So, for the task of hate speech detection model, we will use the Twitter tweets to identify tweets containing Hate Speech.

Dataset & Preprocessing

- We have used Twitter tweets and this datasets contains 31962 instances and 3 features id, tweets and the label.
- We removed the id feature since it does not contribute towards our prediction.
- We removed the stopwords, punctuation marks, urls etc.
- We performed tokenization and lemmatization on our data.
- The dataset had 29720 hate speech tweets and 2242 free speech tweets in the target feature. So we did oversampling of the data for balancing the data.

Models

We implemented various Machine Learning algorithms such as Logistic Regression, Random Forest, Super Vector Machine and Naive Bayes.

We implemented a Neural Network and LSTM model without using transformers and then we built an LSTM model using the Glove Transformer.

Logistic Regression Model with Imbalanced Data

	precision	recall	f1-score	support
0	0.96	1.00	0.98	2984
1	0.95	0.35	0.51	213
accuracy			0.96	3197
macro avg	0.95	0.67	0.74	3197
weighted avg	0.95	0.96	0.95	3197
Logistic Regression, Accuracy Score: 0.955270566155771				

Naive Bayes Model with Imbalanced Data

```
print(classification_report(y_test_tf_wob, y_pred_tf_wob))
```

	precision	recall	f1-score	support
0	0.97	0.86	0.91	5928
1	0.25	0.60	0.35	465
accuracy			0.84	6393
macro avg	0.61	0.73	0.63	6393
weighted avg	0.91	0.84	0.87	6393

Logistic Regression Model with Balanced Data

	precision	recall	f1-score	support
0	0.99	0.95	0.97	2948
1	0.95	0.99	0.97	2996
accuracy			0.97	5944
macro avg	0.97	0.97	0.97	5944
weighted avg	0.97	0.97	0.97	5944
Logistic Regression, Accuracy Score: 0.9703903095558546				

Naive Bayes Model with Balanced Data

	precision	recall	f1-score	support
0	0.84	0.99	0.91	2948
1	0.98	0.81	0.89	2996
accuracy			0.90	5944
macro avg	0.91	0.90	0.90	5944
weighted avg	0.91	0.90	0.90	5944
Naive Bayes, Accuracy Score: 0.8977119784656796				

SVM Model with Balanced Data

	precision	recall	f1-score	support
0	1.00	0.97	0.98	2948
1	0.97	1.00	0.98	2996
accuracy			0.98	5944
macro avg	0.98	0.98	0.98	5944
weighted avg	0.98	0.98	0.98	5944

SVM, Accuracy Score: 0.9835127860026918

Random Forest Model with Balanced Data

	precision	recall	f1-score	support
0	1.00	0.99	0.99	2948
1	0.99	1.00	0.99	2996
accuracy			0.99	5944
macro avg	0.99	0.99	0.99	5944
weighted avg	0.99	0.99	0.99	5944

Random Forest, Accuracy Score: 0.9944481830417228

LSTM Model without Transformer

```
▶ embedding_mat_col=512

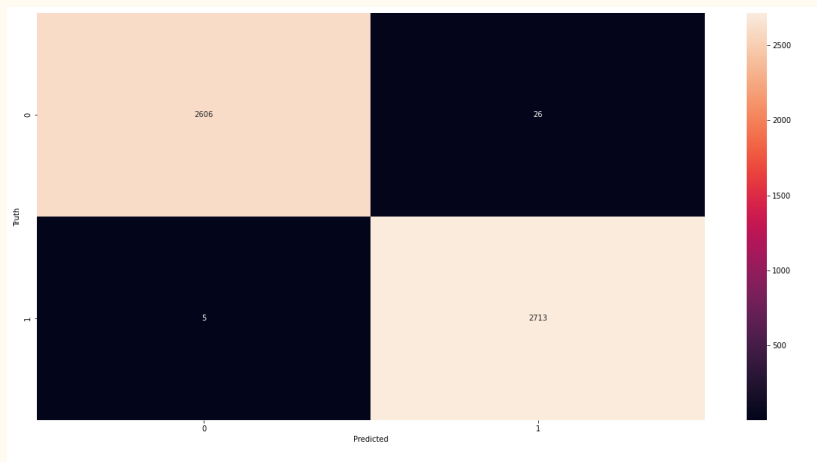
# model = Sequential()
# model.add(Embedding(input_dim=max_words, output_dim=embedding_mat_col, input_length=max_len))
# model.add(SpatialDropout1D(0.4))
# model.add(LSTM(50, dropout=0.4, recurrent_dropout=0.4, input_shape=(None, 512)))
# model.add(Dense(1, activation='softmax'))

# dropout=0.4, recurrent_dropout=0.4,

model_wb = tf.keras.models.Sequential([
    tf.keras.layers.Embedding(input_dim=max_words, output_dim=embedding_mat_col, input_length=max_len),
    tf.keras.layers.LSTM(20, input_shape=(None, 512)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dense(8, activation='relu'),
    tf.keras.layers.Dense(4, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model_wb.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model_wb.summary())
```

Confusion Matrix



```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	2632
1	0.99	1.00	0.99	2718
accuracy			0.99	5350
macro avg	0.99	0.99	0.99	5350
weighted avg	0.99	0.99	0.99	5350

Results using Glove Transformer

```
1 #word embedding - Glove
2
3 from numpy import array
4 from numpy import asarray
5 embeddings_index = dict()
6 f = open('/content/drive/My Drive/glove.6B.100d.txt')
7 for line in f:
8     values = line.split()
9     word = values[0]
10    coefs = asarray(values[1:], dtype='float32')
11    embeddings_index[word] = coefs
12 f.close()
13 print('Loaded %s word vectors.' % len(embeddings_index))
```

NN Model using Glove

```
1 # define model
2 from keras.models import Sequential
3 from keras.layers import Dense
4 from keras.layers import Flatten
5 from keras.layers import Embedding
6 model = Sequential()
7 e = Embedding(vocab_size, 100, weights=[embedding_matrix], input_length=100, trainable=False)
8 model.add(e)
9 model.add(Flatten())
10 model.add(Dense(1, activation='sigmoid'))
```

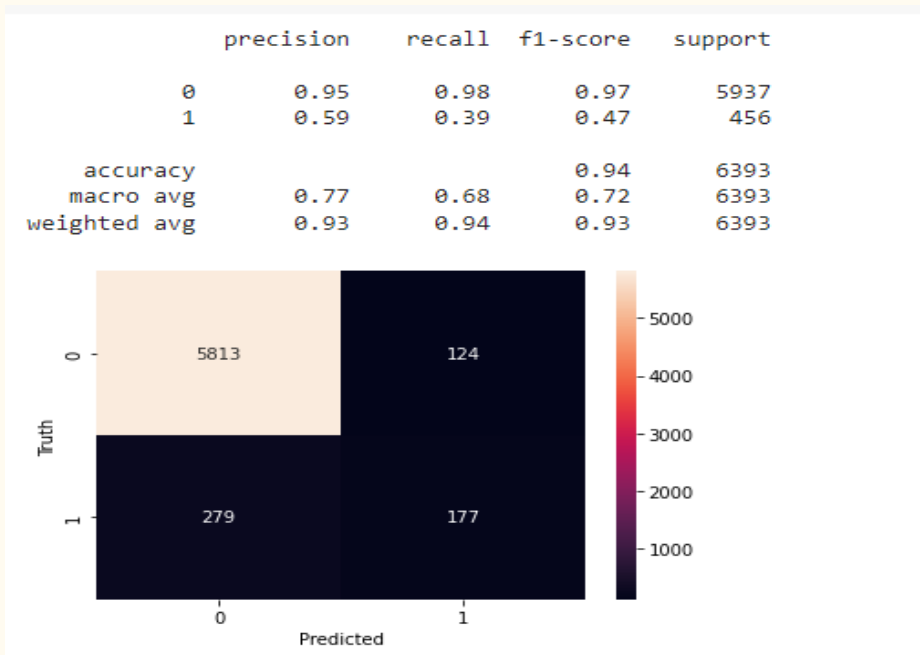
```
1 # compile the model
2 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
3 # summarize the model
4 print(model.summary())
5
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 100)	3384800
flatten (Flatten)	(None, 10000)	0
dense (Dense)	(None, 1)	10001
Total params: 3,394,801		
Trainable params: 10,001		
Non-trainable params: 3,384,800		

None

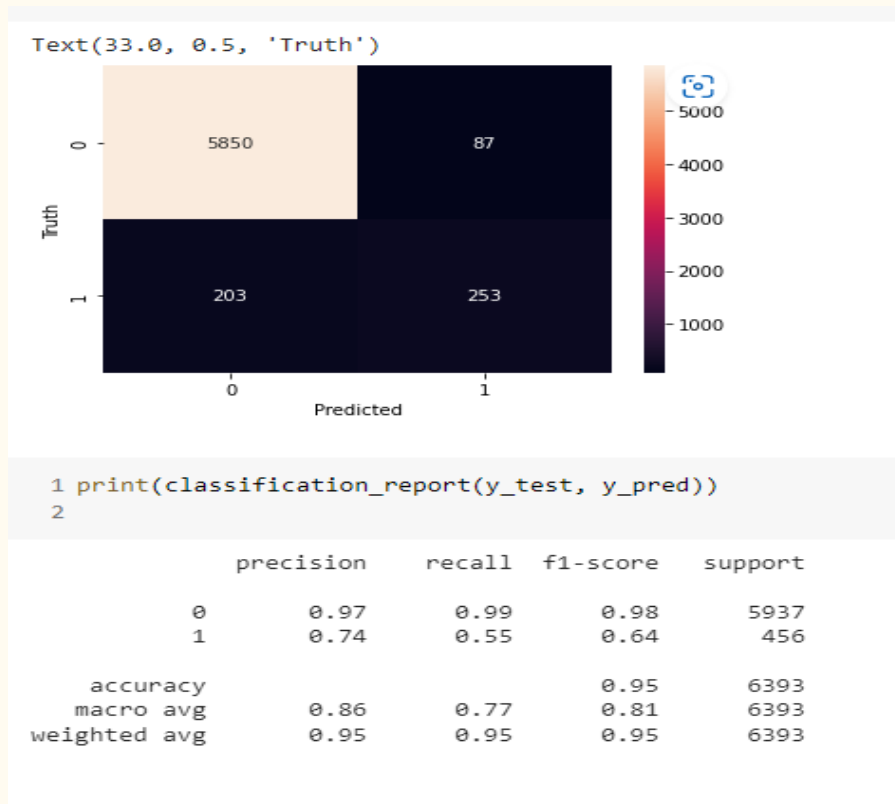
Confusion Matrix and Classification Report



LSTM Model using Glove:

```
1 #LSTM Model
2
3 from keras.layers import Dense, Embedding, LSTM, Bidirectional
4 from keras.models import Sequential
5 from keras.layers import Flatten
6
7 embedding_layer = Embedding(vocab_size, max_length, weights=[embedding_matrix], input_length=max_length, trainable=False)
8
9 embedding_dim = 16
10 input_length = 100
11 model_lstm = Sequential([embedding_layer,
12                           Bidirectional(LSTM(embedding_dim, return_sequences=True)),
13                           Bidirectional(LSTM(embedding_dim, )),
14                           Dense(6, activation='relu'),
15                           Dense(1, activation = 'sigmoid')
16                           ])
17
18 model_lstm.compile(loss = 'binary_crossentropy', optimizer='adam', metrics = ['accuracy'])
```

Confusion Matrix & Classification Report



Conclusion

- The model performed much better when the data is balanced. Hence, we understood the importance of balanced data while training the model.
- LSTM Model performed much better than Naive Bayes and Neural Network Models.